

NatCat: Weakly Supervised Text Classification with Naturally Annotated Resource

Anonymous authors

Abstract

We construct NATCAT, a large-scale resource with rich world knowledge from three online resources: Wikipedia, Reddit, and Stack Exchange. NATCAT consists of document-category pairs derived from manual curation that occurs naturally by their communities. To demonstrate its usefulness, we build general purpose text classifiers by training on NATCAT and evaluate them on a suite of 11 text classification tasks (CATEVAL). We benchmark different modeling choices and resource combinations and show how each task benefits from different NATCAT training resources.¹

1. Introduction

Websites with community contributed content contains ample world knowledge. We seek to improve text classification by constructing large-scale resource of knowledge from freely available online text. NATCAT is a knowledge base automatically constructed from three data sources: Wikipedia, Stackoverflow ², and Reddit ³. With rich world knowledge, we show that NATCAT is a useful resource for zero-shot text classifications.

To demonstrate the usefulness of NATCAT, We train models with it to compute similarity between any document-category pair. As a result, we get weakly supervised (aka. “dataless”) text classifiers that can be used as off-the-shelf classifiers that produce interpretable and relevant topics for any document. They can also be effortlessly ported to a specific topic classification task and categorize documents under the labels of the task. Table 1 illustrates the use of our model on a document from AGNEWS.

To evaluate, we propose CATEVAL, a standardized benchmark for evaluating weakly supervised text classification with a choice of datasets, label descriptions for each dataset, and baseline results. CATEVAL comprises a diverse choice of 11 text classification tasks including both topic-related and sentiment related labels, and contains both single and

1. NATCAT will be available at xxx.yyy/zzz

2. <https://stackoverflow.com/>

3. <https://www.reddit.com/>

Text	Israeli ambassador calls peace conference idea ‘counterproductive’. A broad international peace conference that has reportedly been suggested by Egypt could be “counterproductive” and shouldn’t be discussed until after ...
NatCat	<i>invasions, diplomats, peace, diplomacy, environmentalism, Egypt, patriotism ...</i>
AGNews	<i>international</i> , <i>sports, science technology, business</i>

Table 1: An instance of weakly supervised topic classification from AGNEWS, the highest scoring categories from NATCAT, and ranked AGNEWS categories (true class in bold).

multi-label classification tasks. We show that NATCAT is a valuable resource to build strong weakly supervised text classifiers for future work, and study the impact of NATCAT data domain and pretrained model choice on particular tasks in CATEVAL.

Previous works focus on restricted domains such as medical text [Mullenbach et al., 2018, Rios and Kavuluru, 2018], leverage additional information such as semantic knowledge graphs [Zhang et al., 2019], or carefully exploit weak supervision such as class keywords [Meng et al., 2019] to achieve satisfactory performance. However, they suffer from a limited scope, a need for nontrivial extra supervision that is difficult to obtain in a large amount, and rather complicated methodologies. The scale of NATCAT is much larger than the small dataset-specific experiments in [Yogatama et al., 2017], and the model trained with NATCAT do not require additional supervision at test time such as seed words as in topic modeling approaches [Li et al., 2016a, Chen et al., 2015] or self-training [Meng et al., 2020]. Unlike generic representations such as Explicit Semantic Analysis (ESA) [Chang et al., 2008, Gabrilovich and Markovitch, 2007], with NATCAT, we build neural models by explicitly training on scoring document-category pairs that faithfully approximates the end goal of text classification.

Instead, we consider a more practical setting: is there a readily available resource that we can use to obtain a simple model that can robustly handle a wide range of open-domain text classification tasks? We analyze the gap between weakly supervised and supervised models and show that the mistakes of our models are reasonable and humanlike, suggesting that the construction process of NATCAT is a promising approach of online data mining for building text classifiers.

2. NatCat

In this section, we describe the creation procedures of the NATCAT resource. NATCAT is constructed from three different data sources: Wikipedia, Stack Exchange, and Reddit. Constructed from such three resources, NATCAT naturally contains a wide range of world knowledge that is useful for topical text classifications. We construct document-category pairs from such textual resource where the document falls into the category.

Wikipedia. Wikipedia documents are annotated with categories by the community contributors. The categories of each Wikipedia document can be found at the bottom of

	Wikipedia	Stack Exchange	Reddit
# categories	1,730,447	156	3,000
# documents	2,800,000	2,138,022	7,393,847
avg. # cats. per doc.	86.9	1	1
mode # cats. per doc.	46	1	1

Table 2: Statistics of training sets sampled from the NATCAT dataset, with three different data sources from Wikipedia, Stack Exchange and Reddit.

the page. We obtained Wikipedia documents from Wikimedia Downloads. Wikipedia page-to-category mappings were generated from Wiki SQL dumps using the “categorylinks” and “page” tables. We removed hidden categories by SQL filtering, which are typically maintenance and tracking categories that are unrelated to the document content. We also removed disambiguation categories. After filtering, there are 5.75M documents with at least one category, and a total of 1.19M unique categories. We preprocessed the Wikipedia articles by removing irrelevant information such as the external links at the end of each article. We then removed Wikipedia documents with fewer than 100 non-stopwords.

Some category names are lengthy and specific, e.g., “Properties of religious function on the National Register of Historic Places in the United States Virgin Islands”. These categories are unlikely to be as useful for end users or downstream applications as shorter and more common categories. Therefore, we consider multiple ways of augmenting the given categories with additional categories.

The first way is to use a heuristic method of breaking long category names into shorter ones. We first use stopwords as separators and keep each part of the non-stopword word sequence as a category name. For each category name of a document, we also run a named entity recognizer [Honnibal and Montani, 2017] to find all named entities in that category name, and add them to the category set of the document. This way we expand the existing category names from Wikipedia. For the example category above, this procedure yields the following categories: “religious function”, “the national register of historic places”, “properties”, “historic places”, “the united states virgin islands”, “properties of religious function on the national register of historic places in the united states virgin islands”, “united states virgin islands”, “national register”.

Our second method of expansion is based on the fact that Wikipedia categories can have parent categories and therefore form a hierarchical structure. When expanding the category set by adding its ancestors, there is a trade-off between specificity/relevance and generality/utility of category names. Using only the categories provided for the article yields a small set of high-precision, specific categories. Adding categories that are one or two edges away in the graph increases the total number of training pairs and targets more general/common categories, but some of them will be less relevant to the article. In NATCAT, we include all categories of documents that are up to two edges away.

Stack Exchange. Stack Exchange is a question answering platform where users post and answer questions as a community. Questions on Stack Exchange fall into 308 subareas,

dataset	# test docs.	# labels	# sents. per doc.	# words per doc.	# words per sent.
AGNEWS	7,600	4	1.3	48.8	36.8
DBPEDIA	70k	14	2.4	58.7	24.4
YAHOO	60k	10	5.7	115.8	20.3
20 NEWS GROUPS	7,532	20	15.9	375.4	
EMOTION	16k	10	1.6	19.5	12.4
SST-2	1,821	2	1.0	19.2	19.1
YELP-2	38k	2	8.4	155.1	18.4
AMAZON-2	400k	2	4.9	95.7	19.5
NYTIMES	10k	100	30.0	688.3	22.9
COMMENT	1,287	28	1.3	13.8	10.5
SITUATION	3,525	12	1.8	44.0	24.7

Table 3: Statistics of CATEVAL datasets.

each area having its own site. We construct the document-category pair dataset by pairing question titles or descriptions with their corresponding subareas. Question titles, descriptions and subareas are available from [Chu et al. \[2020\]](#). Many Stack Exchange subareas have their own corresponding “meta” sites. A “Meta” site is meant to discuss the website itself regarding its policy, community, and bugs, etc. When creating this dataset, we merge the subareas with their corresponding “meta” area. This gives us over 2 million documents with 156 categories.

Reddit. Inspired by [Puri and Catanzaro \[2019\]](#), we construct a category classification dataset from Reddit. In our dataset, we propose to classify Reddit post titles to their corresponding subreddit names. We use the OpenWebText⁴ toolkit to get Reddit posts with more than 3 karma and their subreddit names. We only keep the top 3k most frequent subreddits as they better capture the common categories that we are interested in. This gives us over 7 million documents with 3k categories.

Constructed by such three data sources, NATCAT covers a wide range of topics and world knowledge. Table 2 summarizes statistics of training sets we sampled from the NATCAT dataset. Note that all documents from Stack Exchange and Reddit only have one associated category, while a document from Wikipedia may have multiple categories describing it.

3. CatEval Tasks

We use the NATCAT resource to build a zero-shot text classifier. The classifier is applied on a variety of text classification tasks. In this section, we introduce CATEVAL, which comprises a diverse choice of 11 text classification tasks including both topic-related and sentiment-related labels, and contains both single and multi-label classification tasks. For single label topic classification, we have AGNEWS,⁵ DBPEDIA [[Lehmann et al., 2015](#)], YAHOO [[Zhang et al., 2015](#)], and 20 NEWS GROUPS [[Lang, 1995](#)].

4. <https://github.com/jcpeterson/openwebtext>

5. https://www.di.unipi.it/~gulli/AG_corpus_of_news_articles.html

For sentiment classification, we use EMOTION [Klinger et al., 2018], SST-2 [Socher et al., 2013], YELP-2, and AMAZON-2 [Zhang et al., 2015]. EMOTION is a fine-grained sentiment classification tasks with labels expressing various emotions, while the other three are binary sentiment classification tasks differentiating positive and negative sentiments.

As for multi-label topical classification, we have NYTIMES [Sandhaus, 2008], COMMENT,⁶ and SITUATION [Mayhew et al., 2019] datasets. The NYTIMES categories have hierarchical structure, but we merely use the category names from the lowest level. We removed newspaper-specific categories that are not topical in nature.⁷ Of the remaining 2295 categories, we only use the 100 most frequent categories in our experiments, and randomly sample 1 million documents for the training set, 10k for a dev set, and 10k as a test set.⁸

Table 3 summarizes the key statistics of each dataset, including the average number of sentences, average number of words, and average sentence length. They cover a broad range of text classification tasks and can serve as a benchmark for text classifiers.

3.1 Relationship to CatEval Tasks

The goal of weakly supervised text classification is to classify documents into any task-specific categories that are not necessarily seen during training. We demonstrate that NATCAT as a knowledge resource can help to achieve this goal. Even though the category distributions of NATCAT do not exactly match the target distribution in a downstream task, generalization is possible if they are sufficiently similar and covers a diverse set of knowledge. This setting is referred to using several different terms, including dataless classification [Chang et al., 2008], transfer learning [Pan and Yang, 2009], distant supervision [Mintz et al., 2009], and weakly-supervised learning [Zhou, 2017].

We build models with NATCAT that are capable of scoring any candidate label for any document. To evaluate, we take the test sets of standard document classification tasks, use the model to score each label from the set of possible labels for that task, and return the label with the highest score. Therefore, for a given document classification task, we need to specify the name of each label. The choice of label names can have a large impact on performance. As in prior work [Chang et al., 2008, Song and Roth, 2014], we manually choose words corresponding to labels in the downstream tasks. Our models and ESA use the same label names which are provided in the appendix.

As NATCAT is a large textual resource with ample categories, almost all labels in the CATEVAL datasets appear in NATCAT except for some conjunction phrases, such as “written work”, “manufacturing operations and logistics”, and “home and garden”. However, there is no guarantee that the labels in NATCAT have the same definition as the labels in the downstream tasks, and in fact we find such divergences to be causes of error, including when measuring human performance on the text classification tasks. Weakly supervised methods (and humans) are more susceptible to semantic imprecision in label names than supervised methods. The reason we describe our method as “weakly supervised” is because it does not require annotated training data with the same labeling schema and from the

6. <https://dataturks.com/projects/zhiiyubupt/comment>

7. *opinion*, *paid death notices*, *front page*, and *op-ed*

8. Train/dev sets are only used for the supervised baselines.

same distribution as the test set, but rather uses freely-available, naturally-annotated document/category pairs as a training resource.

4. Experiments

4.1 Models and Training

We train BERT [Devlin et al., 2019] and RoBERTa [Liu et al., 2019] on the NATCAT resource for weakly supervised classification. In our experiments, we use BERT-base-uncased (110M parameters) and RoBERTa-base (110M parameters). We formulate NATCAT training as a binary classification task to predict whether a category correctly describes a document. For each document-category pair, we randomly sample 7 negative categories for training. As documents from Wikipedia have multiple positive categories, we randomly sample one positive category for each of them.

We concatenate the category with the document as the input for BERT and RoBERTa: “[CLS] category [SEP] document [SEP]”. In our experiments, we truncate the document to ensure the category-document pair is within 128 tokens.

For BERT and RoBERTa models shown in Table 4, we train models on the whole NATCAT dataset and also the data from a single resource (Wikipedia, Stack Exchange, or Reddit). For each single domain, we train the model for one epoch on 100k instances. For NATCAT combining three domains, we train on 300k instances. The learning rate is set to be 0.00002, and we perform learning rate warmup for 10% of the training steps and then linearly decay the learning rate. As BERT and RoBERTa models are known to suffer from randomness among different runs, we perform each single experiment 5 times under different random seeds and report the median of such five runs. We also do supervised training on EMOTION, NYTIMES, COMMENT, SITUATION with the RoBERTa model. We follow the same training procedure as we train on NATCAT to solve a document-category binary classification task. Our training code is built on Huggingface Transformers [Wolf et al., 2019] and will be released upon publication.

We compare NATCAT trained models to ESA, for which we use their provided code.⁹ We followed the methods of dataless classification from Chang et al. [2008]. Instead of setting a threshold on the number of concepts as in prior work, we use all Wikipedia concepts as we find this improves ESA’s performance.

In preliminary experiments, we experimented with other unsupervised text representation learning approaches, e.g., encode the document and category using pretrained models or pretrained word embeddings (BERT, ELMo, GloVe, etc.), then use cosine similarity as the scoring function for a document-category pair. We also experimented with the pretrained GPT2 model and fine-tuning it on NATCAT [Radford et al., 2019]. However, we found these methods do not perform as well as weakly supervised approaches such as ESA and our approach, so we do not report the results of such methods in this paper.

4.2 Evaluation

We report classification accuracy for all single label classification tasks, including topical and sentiment tasks. For multi-label classification tasks, we use label ranking average

9. github.com/CogComp/cogcomp-nlp/tree/master/dataless-classifier

	Topical (Acc)					Sentiment (Acc)					Multi label (LRAP)				
	ag	dbp	yah	20n	avg	emo	sst	yel	amz	avg	nyt	com	sit	avg	all
NATCAT Trained Models (BERT, RoBERTa, and ensembles)															
B	75.6	82.8	54.9	39.3	63.3	16.1	62.7	70.4	63.6	53.8	49.6	22.6	50.5	41.0	53.3
e	75.4	83.0	55.2	41.7	63.8	16.6	65.7	67.6	68.4	54.6	50.8	22.6	50.8	41.4	54.3
R	68.8	81.9	57.8	36.8	61.3	21.2	65.0	67.3	66.8	55.8	47.7	21.5	52.3	40.5	53.4
e	68.4	85.0	58.5	37.6	62.4	22.3	68.7	75.2	72.4	59.7	49.0	22.1	52.6	41.2	55.6
Other weakly supervised models (ESA, [Yin et al., 2018])															
E	71.2	62.5	29.7	25.1	47.1	9.5	52.1	51.1	51.9	41.2	10.9	22.5	55.6	29.7	40.2
Fully supervised and human performances															
S	92.4	98.7	71.2	85.5	87.0	34.5	97.1	95.6	95.1	80.6	72.5	64.7	75.2	70.8	76.4
H	83.8	88.2	75.0	-	-	-	-	-	-	-	-	-	-	-	-

Table 4: Results of BERT (B) and RoBERTa (R) trained on NATCAT and evaluated on CATEVAL. The metrics are accuracy for single label classification tasks, and LRAP for multi label classification tasks. “e” are ensembles over NATCAT-trained models with 5 random seeds. We compare with the dataless classifier ESA[Chang et al., 2008] (E). We also compare with results from supervised methods (S). All is the average over 11 CATEVAL tasks.

precision (LRAP), a multi-label generalization of the mean reciprocal rank. For each example $i = 1 \dots N$, let $\mathcal{Y}^{(i)} \subseteq \{1 \dots m\} = \mathcal{Y}$ denote the set of gold labels and $f^{(i)} \in \mathbb{R}^m$ denote the model label scores. LRAP is defined as

$$\text{LRAP} \left(\left\{ \mathcal{Y}^{(i)}, f^{(i)} \right\}_{i=1}^N \right) = \frac{1}{N} \sum_{i=1}^N \frac{1}{|\mathcal{Y}^{(i)}|} \sum_{y_{\text{gold}} \in \mathcal{Y}^{(i)}} \frac{\left| \left\{ z_{\text{gold}} \in \mathcal{Y}^{(i)} : f_{z_{\text{gold}}}^{(i)} \geq f_{y_{\text{gold}}}^{(i)} \right\} \right|}{\left| \left\{ y \in \mathcal{Y} : f_y^{(i)} \geq f_{y_{\text{gold}}}^{(i)} \right\} \right|}$$

which achieves the highest value of 1 iff all gold labels are ranked at the top. To directly compare with Yin et al. [2018] in multi label classification tasks, we also use label-wise weighted F1 score in some cases.

4.3 Primary Results

Table 4 summarizes the experimental results of BERT and RoBERTa models trained on NATCAT and evaluated on CATEVAL. RoBERTa trained on NATCAT performs the best on average across tasks, but there are some differences between BERT and RoBERTa. BERT is better on AGNEWS and NYTIMES, both of which are in the newswire domain, as well as 20NG, which also involves some news- or technical-related material. RoBERTa is better on YAHOO as well as better on average in the emotion, binary sentiment, and situation tasks. This may be due to RoBERTa’s greater diversity of training data (web text) compared to BERT’s use of Wikipedia and books.

To provide perspective on the difficulty of the weakly supervised setting, we obtained annotations from 3 human annotators involved in this research project on 60 instances from AGNEWS, 50 from DBPEDIA, and 100 from YAHOO. We showed annotators instances

	Topical (Acc)					Sentiment (Acc)					Multi label (LRAP)				
	ag	dbp	yah	20n	avg	emo	sst	yel	amz	avg	nyt	com	sit	avg	all
BERT models															
W	72.3	86.0	49.0	33.3	60.5	21.3	63.8	64.5	67.0	66.6	41.8	24.3	51.1	39.0	53.0
S	69.0	76.0	59.1	51.2	64.0	18.7	60.1	57.8	57.0	59.1	36.5	24.1	49.9	36.8	51.0
R	70.3	72.8	51.8	49.2	61.7	12.5	61.2	67.0	66.2	65.2	49.8	22.6	52.4	41.5	52.1
N	75.6	82.8	54.9	39.3	63.3	16.1	62.7	70.4	63.6	53.8	49.6	22.6	50.5	41.0	53.3
RoBERTa models															
W	71.7	87.1	53.1	38.8	62.6	22.6	57.2	66.3	69.7	65.3	37.9	23.1	49.9	37.1	52.7
S	65.9	75.5	59.3	19.6	54.7	21.7	59.9	66.2	60.8	62.4	37.7	24.6	47.9	36.8	49.4
R	61.7	71.2	54.0	10.4	49.5	21.3	59.5	57.2	62.9	61.1	42.4	20.6	48.4	37.1	47.1
N	68.8	81.9	57.8	36.8	61.3	21.2	65.0	67.3	66.8	55.8	47.7	21.5	52.3	40.5	53.4

Table 5: Results of BERT and RoBERTa trained on different NATCAT resources (W: Wiki., S: StackEx., R: Reddit, N: NATCAT) and evaluated on CATEVAL. The metrics are accuracy for single label tasks, and LRAP for multi label classification tasks.

and the set of class labels and asked them to choose a single category using their own interpretation and judgment without the ability to look at any training examples. Average accuracies from three annotators on these tasks are reported in Table 4.

We also compare with results from supervised methods. The results of AGNEWS, DBPEDIA, YAHOO, YELP-2 and AMAZON-2 are from Zhang et al. [2015]. The SST-2 result is from Wang et al. [2019]. The 20 NEWS GROUPS result is from Pappagari et al. [2019]. NYTIMES, SITUATION, COMMENT and EMOTION results are fine-tuned RoBERTa models.

In some tasks (AGNEWS and DBPEDIA), supervised models outperform human annotators. We believe this is caused by semantic drift between human interpretation and the actual meaning of the labels as determined in the dataset. Supervised models are capable of learning such nuance from the training data, while an annotator without training is not capable of classifying documents in that way. Weakly supervised models are like human annotators in that they are only capable of classifying documents with the general knowledge they have learned (in this case from large scale naturally-annotated document-category resources).

Yin et al. [2018] builds weakly supervised text classifiers from Wikipedia. Comparing with their reported F1 results in a zero-shot setting (Yahoo: 52.1; Emotion: 21.2; Situation: 27.7), NATCAT trained BERT models yield better results (Yahoo: 54.9; Emotion: 16.1; Situation: 37.1) in topic classification tasks, proving that NATCAT is a valuable resource of topical knowledge.

5. Analysis

5.1 Training Resources

Table 5 shows the model performances when trained on different resources of NATCAT. For each single resource (Wikipedia, StackEx., or Reddit), we train the model for one epoch on 100k instances. We follow the exact same training procedure as described in Subsection 4.1.

All data from three different resources are good at some particular topical classification tasks, most of which can be explained by domain similarities. For example, models trained on Wikipedia are good at DBPEDIA, which can be explained by the fact that DBPEDIA is also built from Wikipedia. Stack Exchange is especially helpful for Yahoo; both are in the domain of community question answering. Models trained on Reddit, which contains a sizable amount of political commentary and news discussion in its most frequent categories, are particularly good at NYTIMES.

Models trained on Stack Exchange do not perform well on most sentiment related tasks. This is likely because Stack Exchange subareas are divided by topic. Wikipedia and Reddit are better resources for training sentiment classifiers, as they cover broader ranges of sentiment and emotion knowledge.

5.2 Training Sizes and Model Variances

While NATCAT has over 10 million documents with over a million categories, we have used a small subset of it due to computational constraints. We compared models trained on 100k and 300k document-category pairs, following the same hyperparameter settings as in Section 4. We find that increasing training size generally harms performance on CATEVAL tasks. For example, averaging over all CATEVAL tasks, BERT trained on Wikipedia is 1.5 points lower when moving from 100k training instances to 300k instances. For Stack Exchange, the gap is 2.1 points. For Reddit, it is 0.7 points.

This is likely due to overfitting on the NATCAT binary classification tasks. As there is a discrepancy between training and evaluation, increasing training data or epochs may not necessarily improve results on downstream tasks. This is a general phenomenon in weakly supervised and zero-shot classification, as we do not have development sets to tune training parameters for such tasks. Similar findings were reported by Puri and Catanzaro [2019], suggesting future work to figure out good ways to do model selection in zero-shot settings.

BERT and RoBERTa are known to suffer from instability in fine-tuning, i.e., training with different random seeds may yield models with vastly different results. We found both models have higher variance on sentiment tasks compared to topic classification. While nontrivial variances are observed, ensembling the 5 models almost always outperforms the median of the individual models.

5.3 Error Analysis

Upon analysis of the confusion matrix of the RoBERTa ensemble predictions on AGNEWS, DBPEDIA, and YAHOO, we observe the following common misclassification instances:

- In AGNEWS, *science & technology* and *international* are often misclassified as *business*.
- In DBPEDIA, *nature* is often misclassified as *animal*, *nature* as *plant*, *written work* as *artist*, and *company* as *transportation*.
- In YAHOO, *society & culture* is often misclassified as *education & reference*, *politics & government*, and *business & finance*. *health* is often misclassified into *science & mathematics*, *family relationships* as *society culture*.

The RoBERTa model trained on NATCAT confuses closely related categories, but it rarely makes mistakes between clearly unrelated concepts. We find that human errors follow

the same pattern: they mostly consist of closely related categories. This suggests that models trained on NATCAT are effective at classifying documents into coarse-grained categories, but fine-grained categorization require annotated training data specific to the task of interest.

6. Related Work

Wikipedia is a classical resource to build weakly-supervised (dataless) text classifiers. A classical work is the dataless text classification approach of Chang et al. [2008], Song and Roth [2014]. This approach uses EXPLICIT SEMANTIC ANALYSIS (ESA) [Gabrilovich and Markovitch, 2007], a method to represent a document and a candidate category as sparse binary indicators of Wikipedia concepts and compute their relatedness by cosine similarity. Wang et al. [2009] learn a universal text classifier based on Wikipedia by extending the dataless approach. Yin et al. [2018] and build models by directly mapping Wikipedia documents to their annotated categories. Puri and Catanzaro [2019] training language models on statements regarding text and corresponding categories from Reddit. NATCAT takes a step further to incorporate text from a variety of freely available text-category pairs from online resources, and show that the diversity of data sources benefit a variety of downstream tasks.

There are other settings of weakly supervised text classification considered in the literature. Some embed text and class labels into the same embedding space and use simple methods for classification [Dauphin et al., 2013, Nam et al., 2016, Li et al., 2016b, Ma et al., 2016]. Others model the presence of a special *unseen* class label and design specialized training and inference procedures to handle it [Shu et al., 2017, Fei and Liu, 2016, Zhang et al., 2019]. Yogatama et al. [2017] report zero-shot text classification experiments with a neural model that jointly embeds words and labels. Mullenbach et al. [2018] and Rios and Kavuluru [2018] focus on medical text. Zhang et al. [2019] models label semantic knowledge in the form of class hierarchies and knowledge graphs. Meng et al. [2019] propose a method by generating pseudo documents based on the keywords and using the documents to train a classifier. NATCAT aims at providing a generic resource of knowledge for text classifiers. Recent works [Zhang et al., 2021, Meng et al., 2020, Shen et al., 2021] use naturally annotated hierarchical tree structure of label taxonomy to improve text classifications.

There is also a wealth of prior work in semi-supervised text classification: using unlabeled text to improve classification performance [Nigam et al., 2000, Howard and Ruder, 2018, Devlin et al., 2019, Liu et al., 2019, Lan et al., 2020, Peters et al., 2018].

7. Conclusion

We proposed NATCAT: a resource of rich world knowledge constructed from online community contributed content. We presented that text classifiers built with NATCAT perform strongly compared to previous approaches of weakly-supervised (dataless) text classifiers. We will release the NATCAT resource, CATEVAL benchmark dataset, our code for running experiments and for evaluation, and our best pretrained model. The NATCAT trained models not only handle any label set but also supply a myriad of interpretable categories for a document off-the-shelf. We believe NATCAT can be a useful resource for applications in natural language processing, information retrieval, and text mining.

References

- Ming-Wei Chang, Lev Ratinov, Dan Roth, and Vivek Srikumar. Importance of semantic representation: Dataless classification. In *AAAI*, 2008.
- Xingyuan Chen, Yunqing Xia, Peng Jin, and John Carroll. Dataless text classification with descriptive LDA. In *AAAI*, 2015.
- Zewei Chu, Mingda Chen, Jing Chen, Miaosen Wang, Kevin Gimpel, Manaal Faruqui, and Xiance Si. How to ask better questions? a large-scale multi-domain dataset for rewriting ill-formed questions. In *AAAI*, 2020.
- Yann N Dauphin, Gokhan Tur, Dilek Hakkani-Tur, and Larry Heck. Zero-shot learning for semantic utterance classification. *arXiv preprint arXiv:1401.0509*, 2013.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL <https://www.aclweb.org/anthology/N19-1423>.
- Geli Fei and Bing Liu. Breaking the closed world assumption in text classification. In *NAACL*, 2016.
- Evgeniy Gabrilovich and Shaul Markovitch. Computing semantic relatedness using wikipedia-based explicit semantic analysis. In *IJCAI*, 2007.
- Matthew Honnibal and Ines Montani. spaCy 2: Natural language understanding with Bloom embeddings, convolutional neural networks and incremental parsing. To appear, 2017.
- Jeremy Howard and Sebastian Ruder. Universal language model fine-tuning for text classification. In *ACL*, 2018.
- Roman Klinger et al. An analysis of annotated corpora for emotion classification in text. In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 2104–2119, 2018.
- Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. ALBERT: A lite BERT for self-supervised learning of language representations. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=H1eA7AEtvS>.
- Ken Lang. Newsweeder: Learning to filter netnews. In *Proceedings of the Twelfth International Conference on Machine Learning*, pages 331–339, 1995.
- Jens Lehmann, Robert Isele, Max Jakob, Anja Jentzsch, Dimitris Kontokostas, Pablo N Mendes, Sebastian Hellmann, Mohamed Morsey, Patrick Van Kleef, Sören Auer, et al. Dbpedia—a large-scale, multilingual knowledge base extracted from wikipedia. *Semantic Web*, 6(2):167–195, 2015.

- Chenliang Li, Jian Xing, Aixin Sun, and Zongyang Ma. Effective document labeling with very few seed words: A topic model approach. In *CIKM*, 2016a.
- Yuezhong Li, Ronghuo Zheng, Tian Tian, Zhiting Hu, Rahul Iyer, and Katia Sycara. Joint embedding of hierarchical categories and entities for concept categorization and dataless classification. *COLING*, 2016b.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- Yukun Ma, Erik Cambria, and Sa Gao. Label embedding for zero-shot fine-grained named entity typing. *COLING*, 2016.
- Stephen Mayhew, Tatiana Tsygankova, Francesca Marini, Zihan Wang, Jane Lee, Xiaodong Yu, Xingyu Fu, Weijia Shi, Zian Zhao, Wenpeng Yin, Karthikeyan K, Jamaal Hay, Michael Shur, Jennifer Sheffield, and Dan Roth. University of Pennsylvania LoReHLT 2019 submission. 2019.
- Yu Meng, Jiaming Shen, Chao Zhang, and Jiawei Han. Weakly-supervised hierarchical text classification. In *AAAI*, 2019.
- Yu Meng, Yunyi Zhang, Jiaxin Huang, Chenyan Xiong, Heng Ji, Chao Zhang, and Jiawei Han. Text classification using label names only: A language model self-training approach. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, 2020.
- Mike Mintz, Steven Bills, Rion Snow, and Dan Jurafsky. Distant supervision for relation extraction without labeled data. In *ACL*, 2009.
- James Mullenbach, Sarah Wiegrefe, Jon Duke, Jimeng Sun, and Jacob Eisenstein. Explainable prediction of medical codes from clinical text. *arXiv preprint arXiv:1802.05695*, 2018.
- Jinseok Nam, Eneldo Loza Mencía, and Johannes Fürnkranz. All-in text: Learning document, label, and word representations jointly. In *Thirtieth AAAI Conference on Artificial Intelligence*, 2016.
- Kamal Nigam, Andrew McCallum, Sebastian Thrun, and Tom Mitchell. Text classification from labeled and unlabeled documents using EM. In *Machine Learning*, 2000.
- Sinno Jialin Pan and Qiang Yang. A survey on transfer learning. *IEEE Transactions on knowledge and data engineering*, 2009.
- Raghavendra Pappagari, Piotr Żelasko, Jesús Villalba, Yishay Carmiel, and Najim Dehak. Hierarchical transformers for long document classification. *arXiv preprint arXiv:1910.10781*, 2019.

- Matthew Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 2227–2237, 2018.
- Raul Puri and Bryan Catanzaro. Zero-shot text classification with generative language models. *arXiv preprint arXiv:1912.10165*, 2019.
- Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. 2019.
- Anthony Rios and Ramakanth Kavuluru. Few-shot and zero-shot multi-label learning for structured label spaces. In *EMNLP*, 2018.
- Evan Sandhaus. The New York Times Annotated Corpus, 2008. URL <https://catalog.ldc.upenn.edu/LDC2008T19>.
- Jiaming Shen, Wenda Qiu, Yu Meng, Jingbo Shang, Xiang Ren, and Jiawei Han. TaxoClass: Hierarchical multi-label text classification using only class names. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4239–4249, Online, June 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.naacl-main.335. URL <https://www.aclweb.org/anthology/2021.naacl-main.335>.
- Lei Shu, Hu Xu, and Bing Liu. Doc: Deep open classification of text documents. In *EMNLP*, 2017.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Y Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1631–1642, 2013.
- Yangqiu Song and Dan Roth. On dataless hierarchical text classification. In *AAAI*, 2014.
- Pu Wang, Charlotta Domeniconi, and Jian Hu. Towards a universal text classifier: Transfer learning using encyclopedic knowledge. In *Data Mining Workshops*, 2009.
- Wei Wang, Bin Bi, Ming Yan, Chen Wu, Zuyi Bao, Liwei Peng, and Luo Si. StructBERT: Incorporating language structures into pre-training for deep language understanding. *arXiv preprint arXiv:1908.04577*, 2019.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, R’emi Louf, Morgan Funtowicz, and Jamie Brew. Hugging-Face’s transformers: State-of-the-art natural language processing. *ArXiv*, abs/1910.03771, 2019.
- Qingyu Yin, Yu Zhang, Weinan Zhang, Ting Liu, and William Yang Wang. Zero pronoun resolution with attention-based neural network. In *Proceedings of the 27th International*

- Conference on Computational Linguistics*, pages 13–23, Santa Fe, New Mexico, USA, August 2018. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/C18-1002>.
- Dani Yogatama, Chris Dyer, Chris Ling, and Phil Blunsom. Generative and discriminative text classification with recurrent neural networks. In *arXiv*, 2017.
- Jingqing Zhang, Piyawat Lertvittayakumjorn, and Yike Guo. Integrating semantic knowledge to tackle zero-shot text classification. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 1031–1040, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1108. URL <https://www.aclweb.org/anthology/N19-1108>.
- Xiang Zhang, Junbo Zhao, and Yann Lecun. Character-level convolutional networks for text classification. In *NIPS*, 2015.
- Yu Zhang, Xiusi Chen, Yu Meng, and Jiawei Han. Hierarchical metadata-aware document categorization under weak supervision. In *WSDM’21*, pages 770–778. ACM, 2021.
- Zhi-Hua Zhou. A brief introduction to weakly supervised learning. *National Science Review*, 5(1):44–53, 2017.

	Topical (Acc)				Sentiment (Acc)			
	AG	DBP	YAH.	20NG	Emo	SST	Yelp	Amz
GPT2 models without candidate answers								
S	55.8	43.7	31.1	25.1	14.1	66.7	66.4	70.0
M	56.4	35.3	32.7	28.1	17.3	66.2	69.5	71.8
L	51.1	42.6	36.7	21.8	17.7	60.4	65.8	69.5
S+NC	51.5	34.2	29.7	14.5	10.2	66.6	68.6	71.1
M+NC	49.9	42.2	28.2	13.0	11.5	53.2	61.5	58.6
GPT2 models with candidate answers								
M	37	7.7	9.9	4.2	5.5	53.9	58.8	60.7
M+NC	72.5	72.6	28.4	4.2	14.9	57.7	60.2	63.7
Puri and Catanzaro [2019]								
1/4	68.3	52.5	52.2	-	-	61.7	58.5	64.5
All	65.5	44.8	49.5	-	-	62.5	74.7	80.2

Table 6: GPT2 results. S/M/L are small/medium/large pretrained GPT2 models, and models with “+NC” fine-tune GPT2 on NATCAT.

Appendix A. Preliminary Experiments with GPT2 Models

We also report preliminary results in adapting GPT2 models to perform CATEVAL tasks. To do so, we construct the following descriptive text: “The document is about [category]: [document content]”, where [category] is replaced by the class label we want to score, and [document content] is the document we want to classify. The descriptive text is tokenized by the BPE tokenizer, truncated to 256 tokens, and fed into the pretrained GPT2 model. The class label with the lowest average loss over all tokens is picked as the predicted label.

The results are shown in the initial rows of Table 6. We find mixed results across tasks, with the GPT2 models performing well on sentiment tasks but struggling on the topical tasks. Increasing GPT2 model size helps in some tasks but hurts in others.

We also fine-tune GPT2 models on NATCAT. Each document in NATCAT is paired with its category to construct the aforementioned descriptive text, and fine-tuned as a language modeling task.¹⁰ The results (upper section of Table 6), are mixed, with the topical accuracies decreasing on average and the sentiment accuracies slightly increasing for GPT2 small but decreasing for GPT2 medium.

A key difference between training GPT2 and BERT/RoBERTa is that with GPT2, we do not explicitly feed information about negative categories. One way to incorporate this information is to construct descriptive text with “candidate categories” following Puri and Catanzaro [2019]¹¹ We sample 7 negative categories and 1 correct category to form the

10. The learning rate (set to 0.00002) follows linear warmup and decay. Following Puri and Catanzaro [2019], we set 1% of training steps as warmup period. We train for one epoch. The maximum sequence length is 256 tokens. We use batch size 8 for GPT2 small (117M parameters) and 2 for GPT2 medium (345M parameters).

11. The descriptive text is as follows: “<|question|> + question + candidate categories + <|endoftext|> + <|text|> + document + <|endoftext|> + <|answer|> + correct category + <|endoftext|>”.

candidates. The results, shown in the middle section of Table 6, improve greatly for some tasks (AG, DBP, and Emo), but drop for the other CATEVAL tasks.

Puri and Catanzaro [2019] also create training data from Reddit. They annotate the text from each outbound weblink with the title of the Reddit post, and the subreddit that the link was posted in. While our Reddit dataset annotates each post title with the name of the subreddit it belongs to.

The GPT2 small model actually outperforms the 1/4-data training setting from Puri and Catanzaro [2019] on the sentiment tasks, though not the All-data training.

Compared to BERT and RoBERTa, it is harder to fine-tune a GPT2 model that performs well across CATEVAL tasks. In fact, there are many ways to convert text classification into language modeling tasks; we explored two and found dramatically different performance from them. It remains an open question how to best formulate text classification for pretrained language models, and how to fine-tune such models on datasets like NATCAT.

Appendix B. Hyperparameters and Model Training

The models and training hyperparameters are described in the main body of the paper. As we run our evaluations in a zero-shot setting, we do not have a development set for parameter tuning and model selection. All of our models (BERT, RoBERTa, GPT2) and training hyperparameters are chosen based on recommended settings from the Huggingface Transformers project [Wolf et al., 2019]. We perform 5 experiment runs with the same hyperparameter settings, except with different random seeds (1, 11, 21, 31, 41), and we report the median performances and standard deviations among different runs.

The following are the sizes of the models we use: BERT-base-uncased (110M), RoBERTa-base (110M), GPT-2 small (117M), GPT-2 medium (345M).

We perform supervised training on EMOTION, NYTIMES, COMMENT and SITUATION. All models are fine-tuned for 3 epochs following the same hyperparameter settings as the main experiment section. We only train on 1000 instances of NYTIMES so we can finish training in a reasonable amount of time (less than 4 hours). COMMENT is trained on the test set as it does not have a training set.

We perform our experiments on single GPUs (including NVIDIA 2080 Ti, Titan X, Titan V, TX Pascal). Training for a single epoch on 300k instances of NATCAT takes about 450 minutes, and for 100k instances it will be around 150 minutes.

For evaluation on CATEVAL, it varies on different tasks. For small tasks like SST-2, it only takes about 20 minutes. For bigger tasks like Amazon-2, it takes about 1200 minutes on a single GPU card for evaluation (we parallelize the evaluation over 40 GPUs so it only takes 30 minutes).

Appendix C. Evaluation Metrics

We use three metrics to evaluate CATEVAL tasks.

Accuracy. Accuracy is the number of correct predictions over the total number of predictions. It is used for single label classification tasks.

	NYT	COMM.	Situ.	AVG
BERT models				
Wiki. 100k	26.4	22.8	23.4	24.6
Wiki. 300k	25.2	22.1	25.7	24.3
StackEx. 100k	28.3	16.0	38.6	28.5
StackEx. 300k	27.9	13.1	37.8	26.5
Reddit 100k	36.4	13.0	32.5	27.7
Reddit 300k	35.4	11.7	28.9	25.4
NATCAT	32.5	16.1	37.1	28.4
RoBERTa models				
Wiki. 100k	25.6	19.1	26.7	23.6
Wiki. 300k	24.4	18.0	29.8	24.0
StackEx. 100k	25.6	8.4	36.2	23.3
StackEx. 300k	23.5	7.6	32.7	21.0
Reddit 100k	36.4	8.5	34.6	24.2
Reddit 300k	35.4	5.5	29.4	21.7
NATCAT	31.0	13.4	36.2	27.2
Other zero-shot				
Yin et al.	-	-	27.7	-

Table 7: F1 scores of multi-label topic classification tasks

Label Ranking Average Precision. This metric is used for evaluating multi-label text classification tasks, and is described in the main body of the paper. We use the scikit-learn implementation from https://scikit-learn.org/stable/modules/generated/sklearn.metrics.label_ranking_average_precision_score.html.

Label Weighted F1 Score. This metric is used to evaluate multi-label text classification tasks. We follow the implementation from Yin et al. [2018].¹²

Table 7 compare the F1 scores of different models on multi label topic classification tasks.

Appendix D. Half Seen Settings

Zero-shot text classification is often defined as training models on some seen labels and testing on an expanded set of both seen and unseen labels [Yogatama et al., 2017, Yin et al., 2018].

We follow the same seen and unseen label splits as Yin et al. [2018], using their v0 and v1 splits. We use the same parameter setting as the main experiments, train BERT models for 3 epochs on the seen label sets, and predict over both seen and unseen labels. We train our models both starting from the original BERT-base-uncased model and the NATCAT-pretrained BERT model. Table 8 summarizes the results (medians over 5 random seeds). The evaluation metrics are accuracy for YAHOO and label-weighted F1 for EMOTION and SITUATION, in order to compare to Yin et al. Pretraining on NATCAT improves BERT’s

¹². https://github.com/yinwenpeng/BenchmarkingZeroShot/blob/master/src/preprocess_situation.py#L204

	Yahoo [Yin et al., 2018]				Emotion				Situation			
	v0		v1		v0		v1		v0		v1	
	seen	unseen	seen	unseen	seen	unseen	seen	unseen	seen	unseen	seen	unseen
Half seen setting												
BERT	73.2	12.9	80.9	9.9	32.8	17.8	34.8	20.3	73.1	50.4	63.9	41.6
+ NATCAT	73.6	16.5	80.1	12.2	33.4	17.6	34.6	19.4	72.8	48.6	61.9	41.3
Yin et al.	<i>72.6</i>	<i>44.3</i>	<i>80.6</i>	<i>34.9</i>	35.6	17.5	37.1	14.2	72.4	48.4	63.8	42.9
NATCAT trained fully unseen setting												
BERT	57.3	52.5	52.5	57.3	18.2	14.6	14.6	18.2	38.4	35.7	35.7	38.4
RoBERTa	63.5	52.1	52.1	63.5	21.4	11.4	11.4	21.4	44.0	30.8	30.8	44.0

Table 8: Experiments of text classification tasks in half seen and fully unseen settings. The evaluation metrics are accuracy for YAHOO and label-weighted F1 for EMOTION and SITUATION, in order to compare to Yin et al. [2018]

results on YAHOO, but it does not show clear improvements on EMOTION and SITUATION in this setting.

Our YAHOO results are not directly comparable to the results from Yin et al. [2018] for several reasons, the most significant being that Yin et al. expand label names using their definitions in WordNet, while we choose to use the plain label names for our experiments.¹³ Another important difference is that Yin et al. [2018] implement a “harsh policy” to impose an advantage to unseen labels by adding an α value to the probabilities of unseen labels. This α value is set by tuning on the development set which contains both seen and unseen labels. However, we do not assume access to a development set with unseen labels.

Table 8 also include the performances of NATCAT trained models on seen and unseens splits of these tasks. On YAHOO, a topical classification task, the model trained on the seen labels perform worse than the model purely trained on NATCAT, making the weakly supervised approach more appealing than this half-seen setting.

Appendix E. Training Sizes

Table 9 shows how the model performances vary with 100k or 300k training instances, following the same hyperparameter settings as in Section 4. We find that increasing training size generally harms performance on CATEVAL tasks. For example, averaging over all CATEVAL tasks, BERT trained on Wikipedia is 1.5 points lower when moving from 100k training instances to 300k instances. For Stack Exchange, the gap is 2.1 points. For Reddit, it is 0.7 points.

This is likely due to overfitting on the NATCAT binary classification tasks. As there is a discrepancy between training and evaluation, increasing training data or epochs may not necessarily improve results on downstream tasks. This is a general phenomenon in weakly supervised and zero-shot classification, as we do not have development sets to tune training

13. Also, Yin et al. formulate the problem as an entailment task, and there are differences in training set sizes.

	Topic	Senti.	Multi-label topic	All
BERT				
Wiki. 100k	60.5	55.3	39.0	53.0
Wiki. 300k	61.9	49.6	38.5	51.5
StackEx. 100k	64.0	49.0	36.8	51.0
StackEx. 300k	62.5	46.7	35.8	48.9
Reddit 100k	61.7	51.6	41.5	52.1
Reddit 300k	62.6	48.8	40.6	51.4
NATCAT 300k	63.3	53.8	41.0	53.3
RoBERTa				
Wiki. 100k	62.6	54.2	37.1	52.7
Wiki. 300k	62.7	55.4	36.9	52.8
StackEx. 100k	54.7	52.1	36.8	49.4
StackEx. 300k	52.8	51.9	35.4	47.4
Reddit 100k	49.5	51.1	37.1	47.1
Reddit 300k	47.9	51.8	37.4	46.2
NATCAT 300k	61.3	55.8	40.5	53.4

Table 9: How training sizes affect model performances on CATEVAL

	Topic	Senti.	Multi-label topic	All
Wiki.	1.1/0.6	3.1/2.8	0.5/0.3	1.3/1.2
StackEx.	0.8/1.2	0.7/2.6	0.5/0.7	0.2/1.2
Reddit	0.8/1.5	3.4/1.8	0.3/1.2	1.4/1.4
NATCAT	0.8/1.2	3.6/1.8	0.7/0.4	1.3/0.2

Table 10: Standard deviations of BERT and RoBERTa model performances on CATEVAL tasks with 5 different random seeds.

parameters for such tasks. Similar findings were reported by [Puri and Catanzaro \[2019\]](#), suggesting future work to figure out good ways to do model selection in zero-shot settings.

Appendix F. Model Variances

BERT and RoBERTa are known to suffer from instability in fine-tuning, i.e., training with different random seeds may yield models with vastly different results. To study this phenomenon in our setting, we performed training in each setting with 5 random seeds and calculate standard deviations for different tasks. As shown in table 10, both models have higher variance on sentiment tasks compared to topic classification. While nontrivial variances are observed, ensembling the 5 models almost always outperforms the median of the individual models.

Appendix G. NatCat Category Names

We list the most frequent 20 categories from Wikipedia, Stack Exchange and Reddit, separated by semicolons.

Wikipedia: years; births by decade; 20th century births; people; people by status; living people; stub categories; works by type and year; works by year; 20th century deaths; establishments by year; 19th century births; years in music; establishments by year and country; establishments by country and year; people by nationality and occupation; 1980s; alumni by university or college in the united states by state; 1970s; 1980s events

Stack Exchange: math; gis; physics; unix; stats; tex; codereview; english; gaming; apple; scifi; drupal; electronics; travel; ell; rpg; meta; mathematica; dba; magento

Reddit: ; AdviceAnimals; politics; worldnews; todayilearned; news; The_Donald; atheism; technology; funny; conspiracy; science; trees; gaming; india; soccer; WTF; reddit.com; Conservative; POLITIC; canada

Appendix H. CatEval Category Names

We list all the category names all tasks in this section, separated by semicolons.

AGNEWS: international; sports; business; science technology

DBPEDIA: company; educational institution; artist; athlete; politician; transportation; building; nature; village; animal; plant; album; film; written work

YAHOO: society culture; science mathematics; health; education reference; computers internet; sports; business finance; entertainment music; family relationships; politics government

20 NEWS GROUPS: atheist christian atheism god islamic; graphics image gif animation tiff; windows dos microsoft ms driver drivers card printer; bus pc motherboard bios board computer dos; mac apple powerbook; window motif xterm sun windows; sale offer shipping forsale sell price brand obo; car ford auto toyota honda nissan bmw; bike motorcycle yamaha; baseball ball hitter; hockey wings espn; encryption key crypto algorithm security; circuit electronics radio signal battery; doctor medical disease medicine patient; space orbit moon earth sky solar; christian god christ church bible jesus; gun fbi guns weapon compound; israel arab jews jewish muslim; gay homosexual sexual; christian morality jesus god religion horus

EMOTION: anger; disgust; fear; guilt; joy; love; no emotion; sadness; shame; surprise

SST-2: Negative; Positive

YELP-2: Negative; Positive

AMAZON-2: Negative; Positive

NYTIMES: new england; real estate; news; britain; theater; new york and region; music theater and dance; your money; russia; iran; art and design; golf; candidates; campaign 2008; new york yankees; israel; pro basketball; healthcare; technology; media entertainment and publishing; family; manufacturing operations and logistics; banking finance and insurance; obituaries; california; media and advertising; health; travel; art; weddings and celebrations; legal; russia and the former soviet union; the city; asia; law enforcement and security; business; week in review; magazine; florida; plays; marketing advertising and pr; new jersey; international; long island; news and features; contributors; texas; style; west; education;

sports; midwest; sunday travel; north america; asia pacific; science; book reviews; united states; westchester; editorials; middle east; markets; south; new york; china; addenda; medicine and health; europe; central and south america; movies; music; road trips; technology telecommunications and internet; washington d.c.; washington; baseball; new york city; arts; books; corrections; iraq; hockey; africa; japan; dance; government philanthropy and ngo; pro football; fashion and style; connecticut; germany; hospitality restaurant and travel; reviews; fashion beauty and fitness; food and wine; letters; usa; france; home and garden; americas; mid atlantic

COMMENT: team; player criticize; audience; sentiment; coach pos; team cav; player praise; team war; game expertise; game observation; refs pos; refs; stats; commercial; player humor; sentiment neg; injury; refs neg; feeling; sentiment pos; coach neg; player; commentary; play; coach; game praise; communication; teasing

SITUATION: water supply; search rescue; evacuation; medical assistance; utilities energy or sanitation; shelter; crime violence; regime change; food supply; terrorism; infrastructure; out of domain